

On the Segmentation of Spoken Documents

A. Researcher and B. Collaborator

Abstract

Reading a document aloud requires deciding where one utterance ends and the next begins. Prior work by Knuth et al. treats the problem as a purely typographic one. We show that layout carries information that punctuation alone does not, particularly around unpunctuated headings. See Fig. 2 for an overview of the pipeline.

1 Introduction

The contributions of this paper are as follows:

- A line-level model of paragraph structure.
- A mapping from characters back to page geometry.
- An evaluation on 400 documents, described in Sec. 4.

The remainder is organised as follows. Section 2 reviews prior work. Section 3 gives the method, and Sec. 4 the evaluation.